

Project 1: File Transport Protocol

Computer Networks (CS-UH 3012) - Fall 2022

1 Code of Conduct

All assignments are graded, meaning we expect you to adhere to the academic integrity standards of NYU Abu Dhabi. To avoid any confusion regarding this, we will briefly state what is and isn't allowed when working on an assignment.

1. Any document and program code that you submit must be fully written by yourself.
2. You can discuss your work with fellow students, as long as these discussions are restricted to general solution techniques. In other words, these discussions should not be about concrete code you are writing, nor about specific results you wish to submit.
3. When discussing an assignment with others, this should never lead to you possessing the complete or partial solution of others, regardless of whether the solution is in paper or digital form, and independent of who made the solution.
4. You are not allowed to possess solutions by someone from a different year or section, by someone from another university, or code from the Internet, etc.
5. There is never a valid reason to share your code with fellow students.
6. There is no valid reason to publish your code online in any form.
7. Every student is responsible for the work they submit. If there is any doubt during the grading about whether a student created the assignment themselves (e.g. if the solution matches that of others), we reserve the option to let the student explain why this is the case. In case doubts remain, or we decide to directly escalate the issue, the suspected violations will be reported to the academic administration according to the policies of NYU Abu Dhabi. More details can be found at:

<https://students.nyuad.nyu.edu/academics/registration/academic-policies/academic-integrity/>

2 Project Objectives

The task of this project is to build a simplified version of the FTP application protocol, consisting of two separate programs: an FTP client and an FTP server. The FTP server is responsible for maintaining FTP sessions and providing file access. The FTP client is split into two components: an FTP user interface and an FTP client to make requests to the FTP server. The client provides a simple user interface with a command prompt asking the user to input a command (see Section 4) that will be issued to the FTP server.

Note that the FTP server must be started first and must support concurrent connections. That is, the server must be able to handle multiple simultaneous requests from different or the same client. This can be done using the `select()`, and `fork()` functions of C. `Fork()` should be used for handling resource-intensive requests/commands with data transfer. For just control connection making and the simple commands, `select()` is sufficient. If you do not use `select()` and with `fork()` for data transfer, there will be a penalty as this is inefficient due to the expensive extra OS resources used by the `fork()` in case it's used for simple commands. This project is a group

project and you must find exactly one partner to work with. Once you have settled on a partner, you need to inform us by sending an email to the TA (smr693@nyu.edu) and the Professor (yz48@nyu.edu). See the due date under submission details and policy. There will be a 10% penalty if the deadline of the group formation is not met.

Please read the complete project description carefully before you start so that you know exactly what is being provided and what functionality you are expected to add. Please also look into the RFC 959 of the FTP protocol.

3 FTP in a Nutshell

In contrast to other protocols that use the same TCP connection for both session control and data transfers, the FTP protocol uses a separate connection for the file transfers and directory listings. That means, an FTP client connects from a random unprivileged port ($N > 1024$) to the server on port 21 which is used as the control channel, i.e. exchanging commands listed in Section 4. For each data transfer the client initiates (using **RETR**, **STOR** or **LIST** command), the client first sends the **PORT N+1** command to the server, specifying what client-side port the server should connect to for the upcoming data transfer. The client then starts listening to port $N+1$ and the server connects from port 20 to the client port to establish the data channel. Once the TCP connection is established, the data transfer is then made through the data channel. The end of the data transfer must be indicated by closing the data channel. Consequently, for every subsequent data transfer (e.g. file request), the client must send the PORT command again to the server to announce a new port number, e.g. $N+2$, $N+3$, etc.. Ports may be reused after some time. Please note that your implementation should only use the active FTP mode. The passive mode does not need to be implemented.

Once the client disconnects (**QUIT** command), the TCP control channel must be closed and the client should terminate.

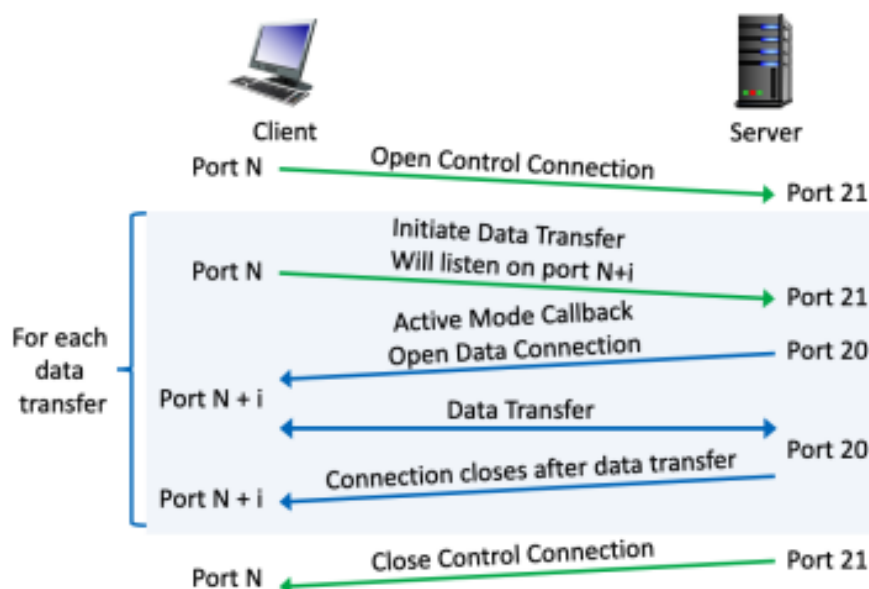


Figure 1: Active FTP Mode: Server initiates data callback connection

4 FTP Commands and Replies

After the establishment of the control TCP connection with the FTPserver, the client process should respond with a prompt *ftp>* waiting for the user's ftp commands. The FTP RFC 959 details a large number of commands that the FTP protocol uses. In this project, we use only a subset of these commands and restrict the transfer mode to stream mode. Note that the commands starting with **!** are executed locally at the client and should execute the corresponding shell commands through the `system()` function.

Below is a list of all the required commands that should be implemented.

PORT h1,h2,h3,h4,p1,p2 This command is used to specify the client IP address and port number for the data channel and is sent automatically by the client before sending a **RETR**, **STOR** or **LIST** command. The server must reply with "200 PORT command successful." before starting the data transfer. The argument is the concatenation of a 32-bit internet host address and a 16-bit TCP port address. This address information is broken into 8-bit fields and the value of each field is transmitted as a decimal number (in character string representation). The fields are separated by commas, where **h1** is the high order 8 bits of the internet client address and **p1** is the high order 8 bits of the client port.

You can use the following formulas to convert the port into character string representation and vice versa:

```
// convert port to p1 and p2
p1 = port/256
p2 = port%256
// convert p1 and p2 to port
port = (p1 * 256) + p2
```

USER *username* This command identifies which user is trying to login to the FTP server. This is mainly used for authentication, since there might be different access controls specified for each user. Once the client issues a **USER** command to the server and the user exists, the server must reply with "331 Username OK, need password.", otherwise with "530 Not logged in."

PASS *password* The client needs to authenticate with the user password by issuing a **PASS** command followed by the user password. The server must reply with either of the following reply codes: "230 User logged in, proceed." or "530 Not logged in.". The user names and passwords should be saved locally in a file called `users.txt` and should be read by the server program upon launch. The file must contain an entry for the username bob with the password donuts.

STOR *filename* This command is used to upload a local file named `filename` from the current client directory to the current server directory.

RETR *filename* This command is used to download a file named *filename* from the current server directory to the current client directory.

LIST This command is used to list all the files under the current server directory.

!LIST This command is used to list all the files under the current client directory.

CWD *foldername* This command is used to change the current server directory. The server must reply with "200 directory changed to pathname/foldername.".

!CWD *foldername* This command is used to change the current client directory.

PWD This command displays the current server directory. The server must reply with "257 pathname.".

!PWD This command displays the current client directory.

QUIT This command quits the FTP session and closes the control TCP connection. The server must reply with "221 Service closing control connection." and the client must terminate after receiving the reply. If this command is issued during an ongoing data transfer, the server/client must abort the data transfer.

Note that the **RETR**, **STOR** and **LIST** commands trigger a data transfer between the client and the server. That is, the server replies with "150 File status okay; about to open data connection." and opens a new TCP data connection to transfer the data. In the control connection, the server creates the socket (on port 21) to which the client binds to open the connection with the server. In the case of data transfer, clients create the sockets on port no: **PORT N + i** (port N is the client side port that it used before to connect to the server), and the server connects to it using bind for opening the data connection with the client. The server connects from port 20 to the client port received (**N + i**) with the **PORT** command and starts transferring the data (**RETR** and **LIST**) or waits to receive a file (**STOR**). If the transferred file already exists, it is replaced. Since this implementation (for simplification purposes) will use the stream transfer mode of FTP, the end of the file is indicated by closing the data connection. That is, the sender (server for **RETR** command, client for **STOR** command) closes the data connection when the file is completely transferred and the server must send a "226 Transfer completed.". Please note that this mode is inherently unreliable, since the client can not determine if the connection closed prematurely or not, i.e. if the file has been entirely transferred or not. You can ignore this unreliability in this implementation. However, since FTP is a stateful protocol, a session must be reset if the client disconnects. Also, since the server must support concurrent file transfers of multiple connected clients and a concurrent file transfer of the same file (or with the same filename) could be a possible scenario, it is highly recommended to save the received file content in a temporary file (e.g. rand#.tmp) or temporary folder during the data transfer and renaming the file to its actual filename after the transfer completed.

Also, the FTP server must be able to respond with appropriate error messages when an error

occurs:

- Any command other than the ones above should not be accepted and should be considered as an invalid FTP command. The server must respond with an "202 Command not implemented."
- The authentication process consists of both a **USER** command and a **PASS** command. No other command should be accepted by the server if the user has not successfully authenticated itself and should result in a "530 Not logged in." response.
- If the user requests an invalid filename or if the directory the user is trying to navigate into does not exist, the client/server must respond with a "550 No such file or directory."
- If the user issues commands in an invalid sequence, the server must respond with a "503 Bad sequence of commands."

Please note, that you don't have to do extensive error checking for the commands, e.g. username or filename containing whitespace, escape characters, etc.

5 Example of an FTP Session

```
220 Service ready for new user.
ftp> USER bob
331 Username OK, need password.
ftp> PASS donuts
230 User logged in, proceed.
ftp> CWD test
200 directory changed to /Users/bob/test
ftp> RETR vanilla_donut.txt
200 PORT command successful.
150 File status okay; about to open. data connection.
226 Transfer completed.
ftp> LIST
200 PORT command successful.
150 File status okay; about to open. data connection.
vanilla_donut.txt
choco_donut.txt
226 Transfer completed.
ftp> QUIT
221 Service closing control connection.
```

6 Testing and Expected Output

Code quality is of particular importance to server robustness in the presence of client errors and malicious attacks. Thus, a large part of this assignment (and programming in general) is knowing how to test, debug and verify your programs. There are many ways to do this; be creative. We would like to know how you tested your FTP server and how you convinced yourself it works as intended. To this end, you should submit your test code along with brief documentation describing what you did to test that your server works. The test cases should include both generic ones that check the server functionality and those that test particular corner cases. In particular, test if files of any type and size are properly transferred and if the server behaves properly if multiple users are connecting/disconnecting.

7 Grading

Description	Score (/20)
Successful compiling of the programs using a Makefile	1
Properly displaying and formatting of the command line user interface	1
Properly handling user authentication	1
Handling concurrent connections and supporting multiple simultaneous users	3
Properly implementing PORT command	2
Properly executing FTP commands and sending correct reply	5
Properly transferring files (binary, plaintext) of any type and size	3
Properly closing the TCP connections (control and data channel)	2
Coding style and usage of meaningful comments	2

8 Submission Details and Policy

Submission Deadlines:

1. Group formation due date: October 6, 2022 (-10% penalty if not met)
2. Submission due date: October 23, 2022 (Midnight)

Submission Format and System: You can directly submit your files as a zip file on Brightspace (<https://brightspace.nyu.edu/>). Due to technical limitations, submissions via email are not accepted.

Late Submissions: Late submissions will be penalized by 10% per 24 hours, with a maximum of 3 days late. In case of a late submission, please upload your zip file to Brightspace and inform the TA and the professor.